

CSE 548: (Design and) Analysis of Algorithms

Divide-and-conquer Algorithms

R. Sekar

1/8

Divide-and-Conquer: A versatile strategy

Steps

- Break a problem into subproblems that are smaller instances of the same problem
- Recursively solve these subproblems
- Combine these answers to obtain the solution to the problem

Benefits

Conceptual simplification

Speed up:

- rapidly (exponentially) reduce problem space
- exploit commonalities in subproblem solutions

Parallelism: Divide-and-conquer algorithms are amenable to parallelization

Locality: Their depth-first nature increases locality, extremely important for today's processors.

2/8

Topics

1. Warmup

Overview
Search
H-Tree

Exponentiation

2. Sorting

Mergesort
Recurrences
Fibonacci
Numbers

Quicksort

Lower Bound
Radix sort

3. Selection

Select k -th min
Priority Queues

4. Closest pair

5. Multiplication

Matrix
Multiplication

Integer

multiplication

6. FFT

Fourier
Transform
DFT
FFT Algorithm
Fast
multiplication

3/8

Binary Search

Problem: Find a key k in an ordered collection

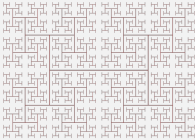
Examples: Sorted array $A[n]$: Compare k with $A[n/2]$, then recursively search in $A[0 \dots (n/2 - 1)]$ (if $k < A[n/2]$) or $A[n/2 \dots n]$ (otherwise)

Binary search tree T : Compare k with $root(T)$, based on the result, recursively search left or right subtree of root.

B-Tree: Hybrid of the above two. Root stores an array M of m keys, and has $m + 1$ children. Use binary search on M to identify which child can contain k , recursively search that subtree.

4/8

H-tree: Planar embedding of full binary tree



Key properties

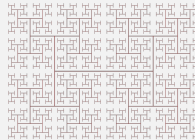
- fractal geometry — divide-and-conquer structure
- n nodes in $O(n)$ area
- all root-to-leaf paths equal in length

Applications

- compact embedding of binary trees in VLSI
- hardware clock distribution

5/8

Divide-and-conquer Construction of H-tree


 $MkHtree(l, b, r, t, n)$
 $horizLine(\frac{b+t}{2}, l + \frac{l+r}{4}, r - \frac{l+r}{4})$
 $vertLine(l + \frac{l+r}{4}, b + \frac{b+t}{4}, t - \frac{b+t}{4})$
 $vertLine(r - \frac{l+r}{4}, b + \frac{b+t}{4}, t - \frac{b+t}{4})$
if $n \leq 4$ **return**
 $MkHtree(l, \frac{b+t}{2}, \frac{l+r}{2}, t, \frac{n}{4})$
 $MkHtree(\frac{l+r}{2}, b, r, t, \frac{n}{4})$
 $MkHtree(l, b, \frac{l+r}{2}, \frac{b+t}{2}, \frac{n}{4})$
 $MkHtree(\frac{l+r}{2}, b, r, \frac{b+t}{2}, \frac{n}{4})$

6/8

Divide-and-conquer Construction of H-tree

Questions

- How compact is the embedding
 - Ratio of minimum distance between nodes and the average area per node
- What is the root-to-leaf path length?
- Can we do better?
- Finally, how can we show that the algorithm is correct?

 $MkHtree(l, b, r, t, n)$
 $horizLine(\frac{b+t}{2}, l + \frac{l+r}{4}, r - \frac{l+r}{4})$
 $vertLine(l + \frac{l+r}{4}, b + \frac{b+t}{4}, t - \frac{b+t}{4})$
 $vertLine(r - \frac{l+r}{4}, b + \frac{b+t}{4}, t - \frac{b+t}{4})$
if $n \leq 4$ **return**
 $MkHtree(l, \frac{b+t}{2}, \frac{l+r}{2}, t, \frac{n}{4})$
 $MkHtree(\frac{l+r}{2}, b, r, t, \frac{n}{4})$
 $MkHtree(l, b, \frac{l+r}{2}, \frac{b+t}{2}, \frac{n}{4})$
 $MkHtree(\frac{l+r}{2}, b, r, \frac{b+t}{2}, \frac{n}{4})$

7/8

Exponentiation

- How many multiplications are required to compute x^n ?
- Can we use a divide-and-conquer approach to make it faster?

 $ExpBySquaring(n, x)$
if $n > 1$
 $y = ExpBySquaring(\lfloor n/2 \rfloor, x^2)$
if $odd(n)$ $y = x * y$
return y
else return x

8/8

Merge Sort

```
function mergesort(a[1...n])
Input: An array of numbers a[1...n]
Output: A sorted version of this array

if n > 1:
    return merge(mergesort(a[1...[n/2]]), mergesort(a[[n/2] + 1...n]))
else:
    return a
```

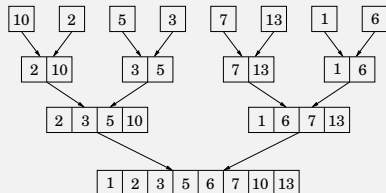
9 / 88

Merge Sort (Continued)

```
function merge(x[1...k], y[1...l])
if k = 0: return y[1...l]
if l = 0: return x[1...k]
if x[1] ≤ y[1]:
    return x[1] ◦ merge(x[2...k], y[1...l])
else:
    return y[1] ◦ merge(x[1...k], y[2...l])
```

10 / 88

Merge Sort Illustration



11 / 88

Merge sort time complexity

- $\text{mergesort}(A)$ makes two recursive invocations of itself, each with an array half the size of A
- $\text{merge}(A, B)$ takes time that is linear in $|A| + |B|$
- Thus, the runtime is given by the recurrence

$$T(n) = 2T\left(\frac{n}{2}\right) + n$$

- In divide-and-conquer algorithms, we often encounter recurrences of the form

$$T(n) = aT\left(\frac{n}{b}\right) + O(n^d)$$

Can we solve them once for all?

12 / 88

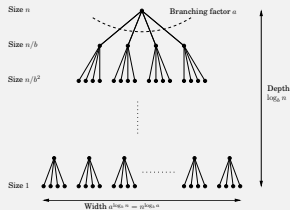
Master Theorem

If $T(n) = aT\left(\frac{n}{b}\right) + O(n^d)$ for constants $a > 0$, $b > 1$, and $d \geq 0$, then

$$T(n) = \begin{cases} O(n^d), & \text{if } d > \log_b a \\ O(n^d \log n) & \text{if } d = \log_b a \\ O(n^{\log_b a}) & \text{if } d < \log_b a \end{cases}$$

13 / 88

Proof of Master Theorem



Can be proved by induction, or by summing up the series where each term represents the work done at one level of this tree.

14 / 88

What if Master Theorem can't be applied?

- Guess and check (prove by induction)
 - expand recursion for a few steps to make a guess
 - in principle, can be applied to any recurrence
- Akra-Bazzi method (not covered in class)
 - recurrences can be much more complex than that of Master theorem

15 / 88

More on time complexity: Fibonacci Numbers

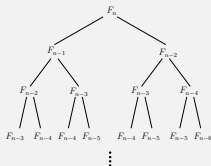
```

function fibl(int n)
    if n = 0 return 0;
    if n = 1 return 1;
    return fibl(n - 1) + fibl(n - 2)
    
```

- Is this algorithm correct? Yes: follows the definition of Fibonacci
- What is its runtime?
 - $T(n) = T(n-1) + T(n-2) + 3$, with $T(k) \leq 2$ for $k < 2$
 - Solution is an exponential function ...
 - Prove this by induction!
- Can we do better?

16 / 88

Structure of calls to *fib*



- Complete binary tree of depth n , contains 2^n calls to *fib*
- But there are only n distinct Fibonacci numbers being computed!
 - Each Fibonacci number computed an exponential number of times!

17/18

Improved Algorithm for Fibonacci

```
function fib2( $n$ )
  int  $f[\max(2, n + 1)]$ ;
   $f[0] = 0$ ;  $f[1] = 1$ ;
  for ( $i = 2$ ;  $i \leq n$ ;  $i++$ )
     $f[i] = f[i - 1] + f[i - 2]$ ;
  return  $f[n]$ 
```

- Linear-time algorithm!
- But wait! We are operating on very large numbers
 - n^{th} Fibonacci number requires approx. $0.694n$ bits
 - Prove *this* by induction!
 - Operation on k -bit numbers require k operations
 - i.e., Computing F_n requires $0.694n \log n$ operations

18/18

Quicksort

```
 $qs(A, l, h)$     /*sorts  $A[l \dots h]$ */
if  $l \geq h$  return;
 $(h, l) =$ 
   $partition(A, l, h)$ ;
 $qs(A, l, h)$ ;
 $qs(A, l, h)$ 

 $partition(A, l, h)$ 
   $k = selectPivot(A, l, h)$ ;  $p = A[k]$ ;
   $swap(A, h, k)$ ;
   $i = l - 1$ ;  $j = h$ ;
  while true do
    do  $i++$  while  $A[i] < p$ ;
    do  $j--$  while  $A[j] > p$ ;
    if  $i \geq j$  break;
     $swap(A, i, j)$ ;
   $swap(A, i, h)$ 
  return  $(j, i + 1)$ 
```

19/18

Analysis of Runtime of *qs*

General case: Given by the recurrence $T(n) = n + T(n_1) + T(n_2)$ where n_1 and n_2 are the sizes of the two sub-arrays after partition.

Best case: $n_1 = n_2 = n/2$. By master theorem, $T(n) = O(n \log n)$

Worst case: $n_1 = 1, n_2 = n - 1$. By master theorem, $T(n) = O(n^2)$

- A fixed choice of pivot index, say, h , leads to worst-case behavior in common cases, e.g., input is sorted.

Lucky/unlucky split: Alternate between best- and worst-case splits.

$$T(n) = n + T(1) + \boxed{T(n-1)} + n \text{ (worst case split)}$$

$$= n + 1 + \boxed{(n-1) + 2T((n-1)/2)} = 2n + 2T((n-1)/2)$$

which has an $O(n \log n)$ solution.

Three-fourths split:

$$T(n) = n + T(0.25n) + T(0.75n) \leq n + 2T(0.75n) = O(n \log n)$$

20/18

Average case analysis of qs

Define input distribution: All permutations equally likely

Simplifying assumption: all elements are distinct. (Nonessential assumption)

Set up the recurrence: When all permutations are equally likely, the selected pivot has an equal chance of ending up at the i^{th} position in the sorted order, for all $1 \leq i \leq n$. Thus, we have the following recurrence for the average case:

$$T(n) = n + \frac{1}{n} \sum_{i=1}^{n-1} (T(i) + T(n-i))$$

Solve recurrence: Cannot apply the master theorem, but since it seems that we get an $O(n \log n)$ bound even in seemingly bad cases, we can try to establish a $cn \log n$ bound via induction.

21/88

Establishing average case of qs

- Establish base case. (Trivial.)
- Induction step involves summation of the form $\sum_{i=1}^{n-1} i \log i$.

Attempt 1: bound $\log i$ above by $\log n$. (Induction fails.)

Attempt 2: split the sum into two parts:

$$\sum_{i=1}^{n/2} i \log i + \sum_{i=n/2+1}^{n-1} i \log i$$

and apply the approx. to each half. (Succeeds with $c \geq 4$.)

Attempt 3: replace the summation with the upper bound

$$\int_{x=1}^n x \log x = \frac{x^2}{2} \left(\log x - \frac{1}{x} \right) \Big|_{x=1}^n$$

(Succeeds with the constraint $c \geq 2$.)

22/88

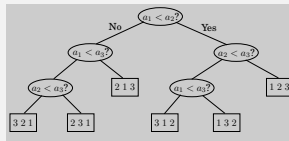
Randomized Quicksort

- Picks a pivot at random
- What is its complexity?
 - For randomized algorithms, we talk about *expected complexity*, which is an average over all possible values of the random variable.
- If pivot index is picked uniformly at random over the interval $[1, h]$, then:
 - every array element is equally likely to be selected as the pivot
 - every partition is equally likely
 - thus, *expected* complexity of *randomized* quicksort is given by the same recurrence as the *average* case of qs .

23/88

Lower bounds for comparison-based sorting

- Sorting algorithms can be depicted as trees: each leaf identifies the input permutation that yields a sorted order.



- The tree has $n!$ leaves, and hence a height of $\log n!$. By Stirling's approximation, $n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n$, so, $\log n! = O(n \log n)$
- No *comparison-based* sorting algorithm can do better!

24/88

Bucket sort

Overview

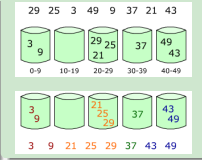
Divide: Partition input into intervals (buckets), based on key values

- Linear scan of input, drop into appropriate bucket

Recurse: Sort each bucket

Combine: Concatenate bin contents

Example



Images from Wikipedia commons

Bucket sort (Continued)

- Bucket sort generalizes quicksort to multiple partitions
 - Combination = concatenation
 - Worst case quadratic bound applies
 - But performance can be much better if input distribution is uniform.
- Exercise:* What is the runtime in this case?
- Used by letter sorting machines in post offices

Counting Sort

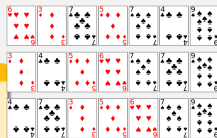
Special case of bucket sort where each bin corresponds to an interval of size 1.

- No need to recurse. Divide = conquered!
- Makes sense only if range of key values is small (usually constant)
- Thus, counting sort can be done in $O(n)$ time!
 - Hmm. How did we beat the $O(n \log n)$ lower bound?*

Radix Sorting

- Treat an integer as a sequence of digits
- Sort digits using counting sort
 - LSD sorting:** Sort first on least significant digit, and most significant digit last. After each round of counting sort, results can be simply concatenated, and given as input to the next stage.
 - MSD sorting:** Sort first on most significant digit, and least significant digit last. Unlike LSD sorting, we cannot concatenate after each stage.
- Note:** Radix sort does not divide inputs into smaller subsets
 - If you think of input as multi-dimensional data, then we break down the problem to each dimension.

- **Stable sorting algorithms:** don't change order of equal elements.
- Merge sort and LSD sort are stable. Quicksort is not stable.



Images from Wikipedia Commons

Why is stability important?

- Effect of sorting on attribute A and then B is the same as sorting on $\langle B, A \rangle$
- LSD sort won't work without this property!
- Other examples: sorting spread sheets or tables on web pages

29 / 88

Sorting strings

- Can use LSD or MSD sorting
 - Easy if all strings are of same length.
 - Requires a bit more care with variable-length strings.
 - Starting point: use a special terminator character $t < a$ for all valid characters a .
- Easy to devise an $O(nl)$ algorithm, where n is the number of strings and l is the maximum size of any string.
 - But such an algorithm is *not* linear in input size.
- **Exercise:** Devise a linear-time string algorithm.
 - Given a set S of strings, your algorithm should sort in $O(|S|)$ time, where

$$|S| = \sum_{s \in S} |s|$$

30 / 88

Obvious approach: Sort, pick k^{th} element — wasteful, $O(n \log n)$

Better approach: Recursive partitioning, search only on one side

$qsel(A, l, h, k)$

```

if  $l = h$  return  $A[l]$ ;
 $(h_1, h_2) = \text{partition}(A, l, h)$ ;
if  $k \leq h_1$ 
    return  $qsel(A, l, h_1, k)$ 
else return  $qsel(A, h_2, h, k)$ 
    
```

Complexity

Best case: Splits are even:

$T(n) = n + T(n/2)$, which has an $O(n)$ solution.

Skewed 10%/90% $T(n) \leq n + T(0.9n)$ — still linear

Worst case: $T(n) = n + T(n-1)$ — **quadratic!**

31 / 88

Intuition: Spend a bit more time to select a pivot that ensures reasonably balanced partitions

MoM Algorithm [Blum, Floyd, Pratt, Rivest and Tarjan 1973]

Time Bounds for Selection

by .

Manuel Blum, Robert W. Floyd, Vaughan Pratt,
Ronald L. Rivest, and Robert E. Tarjan

Abstract

The number of comparisons required to select the i -th smallest of n numbers is shown to be at most a linear function of n by analysis of a new selection algorithm -- PICK. Specifically, no more than $9.595n$ comparisons are ever required. This bound is improved for

32 / 88

$O(n)$ Selection: MoM Algorithm

- Quick select (*qsel*) takes no time to pick a pivot, but then spends $O(n)$ to partition.
- Can we spend more time upfront to make a better selection of the pivot, so that we can avoid highly skewed splits?

Key Idea

- Use the selection algorithm itself to choose the pivot.
 - Divide into sets of 5 elements
 - Compute median of each set ($O(5)$, i.e., constant time)
 - Use selection recursively on these $n/5$ elements to pick their median
 - i.e., choose the median of medians (MoM) as the pivot
- Partition using MoM, and recurse to find k th largest element.

33 / 88

$O(n)$ Selection: MoM Algorithm

Theorem: MoM-based split won't be worse than 30%/70%

Result: Guaranteed linear-time algorithm!

Caveat: The constant factor is non-negligible; use as fall-back if random selection repeatedly yields unbalanced splits.

34 / 88

Selecting maximum element: Priority Queues

Heap

- A tree-based data structure for priority queues

Heap property: H is a heap if for every subtree h of H

$$\forall k \in \text{keys}(h) \quad \text{root}(h) \geq k$$

where $\text{keys}(h)$ includes all keys appearing within h

Note: No ordering of siblings or cousins

- Supports *insert*, *deleteMax* and *max*.
- Typically implemented using arrays, i.e., without an explicit tree data structure



Task of maintaining max is distributed to subsets of the entire set; alternatively, it can be thought of as maintaining several parallel queues with a single head.

Images from Wikimedia Commons

35 / 88

Binary heap

Array representation: Store heap elements in breadth-first order in the array. Node i 's children are at indices $2 * i$ and $2 * i + 1$

- Conceptually, we are dealing with a balanced binary tree

Max: Element at the root of the array, extracted in $O(1)$ time

DeleteMax: Overwrite root with last element of heap. Fix heap – takes $O(\log n)$ time, since only the ancestors of the last node need to be fixed up.

Insert: Append element to the end of array, fix up heap

MkHeap: Fix up the entire heap. Takes $O(n)$ time.

Heapsort: $O(n \log n)$ algorithm, *MkHeap* followed by n calls to *DeleteMax*

36 / 88

Finding closest pair of points

Problem: Given a set of n points in a d -dimensional space, identify the two that have the smallest Euclidean distance between them.



Applications: A central problem in graphics, vision, air-traffic control, navigation, molecular modeling, and so on.

Images from [Wikipedia Commons](#)

37 / 88

Divide-and-conquer closest pair (2D)

Divide: Identify k such that the line $x = k$ divides the points evenly. (Median computation, takes $O(n)$ time.)

Recursive case: Find closest pair in each half.

Combine:

- Can't just take the min of the closest pairs from two halves.
- Need to consider pairs across the divide line — seems that this will take $O(n^2)$ time!

38 / 88

Speeding up search for cross-region pairs

Observation (Key Observation 1)

- Let δ_1 and δ_2 be the minimum distances in each half.
- Need only consider points within $\delta = \min(\delta_1, \delta_2)$ from the dividing line
- We expect that only a small number of points will be within such a narrow strip.
- But in the worst case, every point could be within the strip!

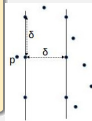
39 / 88

Sparsity condition

Consider a point p on the left δ -strip. How many points q_1, \dots, q_r on the right δ -strip could be within δ from p ?

Observation (Key Observation 2)

- q_1, \dots, q_r should all be within a rectangular $2\delta \times \delta$ as shown
- r can't be too large: q_1, \dots, q_r will crowd together, closer than δ
- **Theorem:** $r \leq 6$



We need to consider at most $6n$ cross-region pairs!

Remains $O(n)$ in higher dimensions as well

40 / 88

Closest pair: Summary

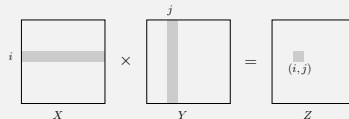
- **Recurrence:** $T(n) = 2T(n/2) + \Omega(n)$, since median computation is already linear-time. Thus, $T(n) = \Omega(n \log n)$.
- To get to $O(n \log n)$, need to
 1. compute the δ -strip in $O(n)$ time
 - Keep the points in each region sorted in x-dimension
 - Takes an additional $O(n \log n)$ time, no change to overall complexity
 2. compute q_1, \dots, q_6 in $O(1)$ time.
 - keep points in each region sorted *also* in y-dimension
 - maintain this order while deleting points outside δ strip
 - in this list, for each p , consider only 12 neighbors — 6 on each side of divide

41/81

Matrix Multiplication

The product Z of two $n \times n$ matrices X and Y is given by

$$Z_{ij} = \sum_{k=1}^n X_{ik} Y_{kj} \quad \text{— leads to an } O(n^3) \text{ algorithm.}$$



42/81

Divide-and-conquer Matrix Multiplication

Divide X and Y into four $n/2 \times n/2$ submatrices

$$X = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \text{ and } Y = \begin{bmatrix} E & F \\ G & H \end{bmatrix}$$

Recursively invoke matrix multiplication on these submatrices:

$$XY = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \begin{bmatrix} E & F \\ G & H \end{bmatrix} = \begin{bmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{bmatrix}$$

Divided, but did not conquer! $T(n) = 8T(n/2) + O(n^2)$, which is still $O(n^3)$

43/81

Strassen's Matrix Multiplication

Strassen showed that 7 multiplications are enough:

$$XY = \begin{bmatrix} P_6 + P_3 + P_4 - P_2 & P_1 + P_2 \\ P_3 + P_4 & P_1 - P_3 + P_5 - P_7 \end{bmatrix} \quad \text{where}$$

$$\begin{aligned} P_1 &= A(F - H) & P_5 &= (A + D)(E + H) \\ P_2 &= (A + B)H & P_6 &= (B - D)(G + H) \\ P_3 &= (C + D)E & P_7 &= (A - C)(E + F) \\ P_4 &= D(G - E) \end{aligned}$$

Now, the recurrence $T(n) = 7T(n/2) + O(n^2)$ has $O(n^{\log_2 7} = n^{2.81})$ solution!

Best-to-date complexity is about $O(n^{2.4})$, but this algorithm is not very practical.

44/81

Karatsuba's Algorithm

Same high-level strategy as Strassen — but predates Strassen.

Divide: n -digit numbers into halves, each with $n/2$ -digits:

$$\begin{aligned} a &= \boxed{a_1} \boxed{a_0} = 2^{n/2} a_1 + a_0 \\ b &= \boxed{b_1} \boxed{b_0} = 2^{n/2} b_1 + b_0 \\ ab &= 2^n a_1 b_1 + 2^{n/2} (a_1 b_0 + b_1 a_0) + a_0 b_0 \end{aligned}$$

Key point — Instead of 4 multiplications, we can get by with 3 since:

$$a_1 b_0 + b_1 a_0 = (a_1 + a_0)(b_1 + b_0) - a_1 b_1 - a_0 b_0$$

Recursively compute $a_1 b_1$, $a_0 b_0$ and $(a_1 + a_0)(b_1 + b_0)$.

Recurrence $T(n) = 3T(n/2) + O(n)$ has an $O(n^{\log_2 3} = n^{1.59})$ solution!

Note: This trick for using 3 (not 4) multiplications noted by Gauss (1777-1855) in the context of complex numbers.

Toom-Cook Multiplication

- Generalize Karatsuba
 - Divide into $n > 2$ parts
- Can be more easily understood when integer multiplication is viewed as a polynomial multiplication.

Integer Multiplication Revisited

- An integer represented using digits

$$a_{n-1} \dots a_0$$

over a base d (i.e., $0 \leq a_i < d$) is very similar to the polynomial

$$A(x) = \sum_{i=0}^{n-1} a_i x^i$$

Specifically, the value of the integer is $A(d)$.

- Integer multiplication follows the same steps as polynomial multiplication:

$$a_{n-1} \dots a_0 \times b_{n-1} \dots b_0 = (A(x) \times B(x))(d)$$

Polynomials: Basic Properties

Horner's rule

An n^{th} degree polynomial $\sum_{i=0}^n a_i x^i$ can be evaluated in $O(n)$ time:

$$((\dots((a_n x + a_{n-1})x + a_{n-2})x + \dots + a_1)x + a_0)$$

Roots and Interpolation

- An n^{th} degree polynomial $A(x)$ has exactly n roots r_1, \dots, r_n . In general, r_i 's are complex and need not be distinct.
- It can be represented as a product of sums using these roots:

$$A(x) = \sum_{i=1}^n a_i x^i = \prod_{i=1}^n (x_i - r_i)$$

- Alternatively, $A(x)$ can be specified uniquely by specifying $n+1$ points (x_i, y_i) on it, i.e., $A(x_i) = y_i$.

Operations on Polynomials

Representation	Add	Mult
Coefficients	$O(n)$	$O(n^2)$
Roots	?	$O(n)$
Points	$O(n)$	$O(n)$

Note: Point representation is the best for computation! But usually, only the coefficients are given

Solution: Convert to point form by *evaluating* $A(x)$ at selected points.

But conversion defeats the purpose: requires $O(n)$ evaluations, each taking $O(n)$ time, thus we are back to $O(n^2)$ total time.

Toom (and FFT) Idea: Choose evaluation points judiciously to speed up evaluation

49/81

Matrix representation of Polynomial Evaluation

Given a polynomial

$$A(x) = \sum_{i=0}^{n-1} a_i x^i$$

choose m points x_0, \dots, x_m for its evaluation.

Evaluation can be expressed using matrix multiplication:

$$\begin{bmatrix} p_0 \\ p_1 \\ p_2 \\ \vdots \\ p_m \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \cdots & x_1^{n-1} \\ 1 & x_2 & x_2^2 & \cdots & x_2^{n-1} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_m & x_m^2 & \cdots & x_m^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

50/81

Multiplication using Point Representation

- Let $A(x)$ and $B(x)$ be polynomials representing two numbers
- Evaluate both polynomials at chosen points x_0, \dots, x_m

$$P = \mathbf{X}A \quad Q = \mathbf{X}B$$

where P, \mathbf{X}, A, Q and B denote matrices as in last page

- Compute point-wise product

$$\begin{bmatrix} r_0 \\ \vdots \\ r_m \end{bmatrix} = \begin{bmatrix} p_0 * q_0 \\ \vdots \\ p_m * q_m \end{bmatrix}$$

- Compute polynomial C corresponding to R

$$R = \mathbf{X}C \Rightarrow C = \mathbf{X}^{-1}R$$

- To avoid overflow, m should be $\text{degree}(A) + \text{degree}(B) + 1$ for R

51/81

Improving complexity ...

- Key problem: Complexity of computing \mathbf{X} and its inverse \mathbf{X}^{-1}
- Toom strategy:
 - Use low-degree polynomials e.g., Toom-2 = Karatsuba uses degree 1.
 - represents an n -bit number as a 2-digit number over a large base $d = 2^{n/2}$
 - Fix evaluation points for a given degree polynomial so that \mathbf{X} and \mathbf{X}^{-1} can be precomputed
 - For Toom-2, $x_0 = 0, x_1 = 1, x_2 = \infty$. (Define $A(\infty) = a_{n-1}$)
 - Choose points so that expensive multiplications can be avoided while computing $P = \mathbf{X}A, Q = \mathbf{X}B$ and $C = \mathbf{X}^{-1}R$
- Toom- N on n -digit numbers needs $2N - 1$ multiplications on n/N digit numbers:

$$T(n) = (2N - 1)T(n/N) + O(n)$$

which, by Master theorem, has a solution $O(n^{\log_N(2N-1)})$ solution

52/81

Karatsuba revisited as Toom-2

Given evaluation points $x_0 = 0, x_1 = 1, x_2 = \infty$,

$$\mathbf{X} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad \mathbf{X}\mathbf{A} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ 0 \end{bmatrix} = \begin{bmatrix} a_0 \\ a_0 + a_1 \\ a_1 \end{bmatrix}$$

Similarly

$$\mathbf{X}\mathbf{B} = \begin{bmatrix} b_0 \\ b_0 + b_1 \\ b_1 \end{bmatrix}$$

Point-wise multiplication yields:

$$\mathbf{R} = \begin{bmatrix} a_0 b_0 \\ (a_0 + a_1)(b_0 + b_1) \\ a_1 b_1 \end{bmatrix}$$

and so on ...

53 / 88

Limitations of Toom

- In principle, complexity can be reduced to $n^{1+\epsilon}$ for arbitrarily small positive ϵ by increasing N
- In reality, the algorithm itself depends on the choice of N . Specifically, constant factors involved increase rapidly with N .
- As a practical matter, $N = 4$ or 5 is where we stop.
- Question: Can we go farther?

54 / 88

FFT and Schonhage-Strassen

- Key idea: evaluate polynomial on the complex plane
- Choose powers of N th complex root of unity as the points for evaluation
- Enables sharing of operations in computing $\mathbf{X}\mathbf{A}$ so that it can be done in $O(N \log N)$ time, rather than $O(N^2)$ time needed for the naive matrix-multiplication based approach

55 / 88

FFT to the Rescue!

Matrix form of DFT and interpretation as polynomial evaluation:

$$\begin{bmatrix} s_0 \\ s_1 \\ \vdots \\ s_j \\ \vdots \\ s_{N-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ 1 & \omega & \omega^2 & \cdots & \omega^{N-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \cdots & \omega^{j(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{N-1} & \omega^{2(N-1)} & \cdots & \omega^{(N-1)(N-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_j \\ \vdots \\ a_{N-1} \end{bmatrix}$$

- **Voila!** FFT computes $A(x)$ at N points ($x_j = \omega^j$) in $O(N \log N)$ time!
- $O(N \log N)$ integer multiplication

Convert to point representation using FFT	$O(N \log N)$
Multiply on point representation	$O(N)$
Convert back to coefficients using FFT ⁻¹	$O(N \log N)$

56 / 88

FFT to the Rescue!

$$\begin{bmatrix} s_0 \\ s_1 \\ \vdots \\ s_j \\ \vdots \\ s_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{j(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_j \\ \vdots \\ a_{n-1} \end{bmatrix}$$

FFT can be thought of as a clever way to choose points:

- Evaluations at many distinct points “collapse” together
- This is why we are left with $2T(n/2)$ work after division, instead of $4T(n/2)$ for a naive choice of points.

57/88

FFT-based multiplication: More careful analysis ...

- Computations on complex or real numbers can lose precision.
 - For integer operations, we should work in some other ring — usually, we choose a ring based on modulo arithmetic.
 - Ex: in mod 33 arithmetic, 2 is the 10th root of 1, i.e., $2^{10} \equiv 1 \pmod{33}$

More generally, 2 is the n th root of unity modulo $(2^{n/2} + 1)$
- Point-wise additions and multiplications are not $O(1)$.
 - We are adding up to n numbers (“digits”) — we need $\Omega(\log n)$ bits
 - So, total cost increases by at least $\log n$, i.e., $O(n \log^2 n)$.
- [Schonhage-Strassen '71] developed $O(n \log n \log \log n)$ algorithm: recursively apply their technique for “inner” operations.

58/88

Integer Multiplication Summary

- Algorithms implemented in libraries for arbitrary precision arithmetic, with applications in public key cryptography, computer algebra systems, etc.
- GNU MP is a popular library, uses various algorithms based on input size: naive, Karatsuba, Toom-3, Toom-4, or Schonhage-Strassen (at about 50K digits).
- Karatsuba is Toom-2. Toom-N is based on
 - Evaluating a polynomial at $2N$ points,
 - performing point-wise multiplication, and
 - interpolating to get back the polynomial, while
 - minimizing the operations needed for interpolation

59/88

Fast Fourier Transformation

One of the most widely used algorithms — yet most people are unaware of its use!

Solving differential equations: Applied to many computational problems in engineering, e.g., heat transfer

Audio: MP3, digital audio processors, music/speech synthesizers, speech recognition, ...

Image and video: JPEG, MPEG, vision, ...

Communication: modulation, filtering, radars, software-defined radios, H.264, ...

Medical diagnostics: MRI, PET, ultrasound, ...

Quantum computing: See text Ch. 10

Other: Optics, data compression, seismology, ...

60/88

Fourier Series



Images from Wikimedia Commons. 61/81

Theorem (Fourier Theorem)

Any (sufficiently smooth) function with a period T can be expressed as a sum of series of sinusoids with periods T/n for integral n .

$$a(t) = \sum_{n=0}^{\infty} (d_n \sin(2\pi nt/T) + e_n \cos(2\pi nt/T))$$

Fourier Series

Example: Touch tone button 1

Using the identity

$$e^{ix} = \cos x + i \sin x$$

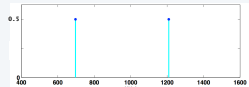
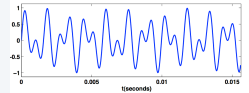
Fourier series becomes

$$a(t) = \sum_{n=-\infty}^{\infty} c_n e^{2\pi i n t / T}$$

It can be shown

$$c_n = \int_0^T a(t) e^{-2\pi i n t / T} dt$$

For real $a(t)$, $c_n = c_{-n}^*$.



Images from Cleve Moler, Numerical computing with MATLAB

62/81

Fourier Transform

- What if a is not periodic?
- Maybe we can start with the Fourier series definition for c_n

$$c_n = \int_0^T a(t) e^{-2\pi i n t / T} dt$$

and let $T \rightarrow \infty$?

- Frequencies are not discrete any more, as the "fundamental frequency" $f = 1/T \rightarrow 0$
- Instead of discrete coefficients c_n we will have a continuous function — call it $s(f)$.
- $\mathcal{F}(a)$ denotes a 's Fourier transform
- \mathcal{F} is almost self-inverting: $\mathcal{F}(\mathcal{F}(a(t))) = a(-t)$

$$s(f) = \int_{-\infty}^{\infty} a(t) e^{-2\pi i f t} dt$$

63/81

How do Fourier Series/Transform help?

Differential equations: Turn non-integrable functions into a sum of easily integrable ones.

Some problems easier to solve in frequency domain:

Filtering: filter out noise, tuning, ...

Compression: eliminate high frequency components, ...

Convolution: Convolution in time domain becomes (simpler) multiplication in frequency domain.

Definition (Convolution)

$$(a * b)(t) = \int_{-\infty}^{\infty} a(t-x)b(x) dx$$

Theorem (Convolution)

$$\mathcal{F}(a * b)(t) = \mathcal{F}(a(t))\mathcal{F}(b(t))$$

64/81

Discrete Fourier Transform

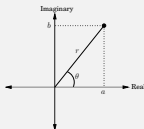
- Real-world signals are typically sampled
 - DFT is a formulation of FT applicable to such samples
- Nyquist rate:** A signal with highest frequency $n/2$ can be losslessly reconstructed from n samples.
- DFT of time domain samples a_0, \dots, a_{n-1} yields frequency domain samples s_0, \dots, s_{n-1} :

$$s_f = \sum_{t=0}^{n-1} a_t e^{-2\pi i f t / n} \quad \text{cf. } s(f) = \int_{-\infty}^{\infty} a(t) e^{-2\pi i f t} dt$$

Note: DFT formulation can be derived from FT by treating the sampling process as a multiplication by a sequence of impulse functions separated by the sampling interval

65 / 81

Background: Complex Plane, Polar Coordinates



The complex plane

$z = a + bi$ is plotted at position (a, b) .

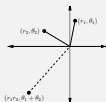
Polar coordinates: rewrite as $z = r(\cos \theta + i \sin \theta) = r e^{i\theta}$, denoted (r, θ) .

- length $r = \sqrt{a^2 + b^2}$.
- angle $\theta \in [0, 2\pi)$: $\cos \theta = a/r, \sin \theta = b/r$.
- θ can always be reduced modulo 2π .

Examples:	Number	-1	i	$5 + 5i$
	Polar coords	$(1, \pi)$	$(1, \pi/2)$	$(5\sqrt{2}, \pi/4)$

66 / 81

Polar Coordinates and Multiplication



Multiply the lengths and add the angles:

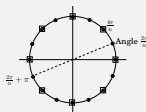
$$(r_1, \theta_1) \times (r_2, \theta_2) = (r_1 r_2, \theta_1 + \theta_2).$$

For any $z = (r, \theta)$,

- $-z = (r, \theta + \pi)$ since $-1 = (1, \pi)$.
- If z is on the **unit circle** (i.e., $r = 1$), then $z^n = (1, n\theta)$.

67 / 81

Roots of unity on Complex Plane



Solutions to the equation $z^n = 1$.

By the multiplication rule: solutions are $z = (1, \theta)$, for θ a multiple of $2\pi/n$ (shown here for $n = 16$).

For even n :

- These numbers are **plus-minus paired**: $-(1, \theta) = (1, \theta + \pi)$.
- Their squares are the $(n/2)$ nd roots of unity, shown here with boxes around them.

68 / 81

Matrix representation of DFT

- Given time domain samples a_t for $t = 0, 1, \dots, n-1$,
- Compute frequency domain samples s_f for $f = 0, 1, \dots, n-1$

$$s_f = \sum_{t=0}^{n-1} a_t e^{-2\pi i f t / n} = \sum_{t=0}^{n-1} a_t \left(e^{-2\pi i / n} \right)^{ft} = \sum_{t=0}^{n-1} a_t \omega^{ft}$$

where $\omega = e^{-2\pi i / n}$ is the n th complex root of unity

$$\begin{bmatrix} s_0 \\ s_1 \\ s_2 \\ \vdots \\ s_j \\ \vdots \\ s_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ 1 & \omega^2 & \omega^4 & \dots & \omega^{2(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{j(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_j \\ \vdots \\ a_{n-1} \end{bmatrix}$$

69 / 88

Matrix representation of DFT

Note that $e^{-2\pi i / n}$ represents the n th root of 1, denoted ω .

$$\begin{bmatrix} s_0 \\ s_1 \\ s_2 \\ \vdots \\ s_j \\ \vdots \\ s_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ 1 & \omega^2 & \omega^4 & \dots & \omega^{2(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{j(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_j \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Possible interpretations of these matrix equations:

- Simultaneous equations that can be solved
- Change of basis (rotate coordinate system)
- Evaluation of polynomial $\sum_{k=0}^{n-1} a_k x^k$ at $x = \omega^j, 0 \leq j < n$.

70 / 88

Speeding up FFT Computation

- Matrix multiplication formulation has an obvious divide-and-conquer implementation

$$M_n \vec{A}_{0 \dots (n-1)} = \begin{bmatrix} M_{n/2}^{11} & M_{n/2}^{12} \\ M_{n/2}^{21} & M_{n/2}^{22} \end{bmatrix} \begin{bmatrix} \vec{A}_{0 \dots [n/2]} \\ \vec{A}_{[n/2] \dots (n-1)} \end{bmatrix}$$

But this algorithm still takes $O(n^2)$ time

- ... but wait! — there are only $O(n)$ distinct elements in the square matrix M_n .
- $O(n)$ repetitions of each element in M_n , so there is significant scope for sharing operations on submatrices!

71 / 88

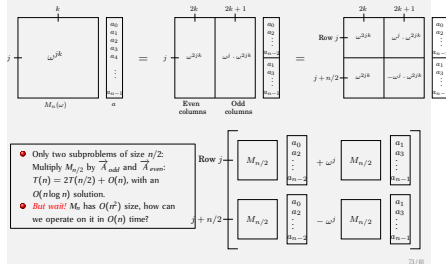
Observations about $M(\omega)$



- Two successive columns differ by a factor ω^j in the j th row
- Rows that are $n/2$ rows apart differ by a factor of $\omega^{kn/2}$ in the k th column
 - Note that $\omega^{n/2} = -1$, so they differ by a factor of -1 on odd columns, and are identical on even columns.

72 / 88

DFT Matrix Multiplication, Rearranged ...



73/88

FFT Algorithm

function FFT(a, ω)

Input: An array $a = (a_0, a_1, \dots, a_{n-1})$, for n a power of 2
 A primitive n th root of unity, ω

Output: $M_n(\omega)a$

if $\omega = 1$: return a

$(s_0, s_1, \dots, s_{n/2-1}) = \text{FFT}(a_0, a_2, \dots, a_{n-2}, \omega^2)$

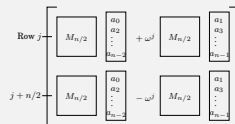
$(s'_0, s'_1, \dots, s'_{n/2-1}) = \text{FFT}(a_1, a_3, \dots, a_{n-1}, \omega^2)$

for $j = 0$ to $n/2 - 1$:

$r_j = s_j + \omega^j s'_j$

$r_{j+n/2} = s_j - \omega^j s'_j$

return $(r_0, r_1, \dots, r_{n-1})$



74/88

Convolution in the Discrete World

$$(\vec{A}_n * \vec{B}_m)_t = \sum_{x=0}^{m-1} a_{t-x} b_x \quad \text{cf. } (a * b)(t) = \int_{-\infty}^{\infty} a(t-x)b(x)dx$$

Linear convolution: $a_{t-x} = 0$ if $x > t$

Circular convolution: $a_{t-x} = a_{t-n+x}$ if $x > n$. (Equivalent to treating A as a periodic function.)

Zero-extended convolution: First extend A and B to have $m + n - 1$ samples by letting $A_y = 0$ for $m \leq y < m + n$ and $B_z = 0$ for $n \leq z < m + n$.

With zero-extension, the definitions of linear and circular conventions match, and hence become equivalent. Hence, we will deal only with zero-extended convolution.

Theorem (Discrete Convolution) $\mathcal{F}(\vec{A}_n * \vec{B}_m) = \mathcal{F}(\vec{A}_n) \mathcal{F}(\vec{B}_m)$

75/88

Why this fascination with convolution?

- Computationally, convolution is a loop to add products
- The convolution theorem says we can replace this $O(n)$ loop by a single operation on the DFT. *That is fascinating!*
 - *Wait a minute!* What about the cost of computing \mathcal{F} first?
- If we use FFT, then we the computation of \mathcal{F} and its inversion will still be $O(n \log n)$, not quadratic.

- Can we use FFT as a building block to speed up algorithms for other problems?
 - Integer multiplication looks like a convolution, and usually takes $O(n^2)$. Can we make it $O(n \log n)$?

76/88

FFT to the Rescue!

Matrix form of DFT and interpretation as polynomial evaluation:

$$\begin{bmatrix} s_0 \\ s_1 \\ \vdots \\ s_j \\ \vdots \\ s_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{j(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_j \\ \vdots \\ a_{n-1} \end{bmatrix}$$

- **Voila!** FFT computes $A(x)$ at n points ($x_i = \omega^i$) in $O(n \log n)$ time!
- $O(n \log n)$ integer multiplication

Convert to point representation using FFT	$O(n \log n)$
Multiply on point representation	$O(n)$
Convert back to coefficients using FFT ⁻¹	$O(n \log n)$

77/88

FFT to the Rescue!

$$\begin{bmatrix} s_0 \\ s_1 \\ \vdots \\ s_j \\ \vdots \\ s_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^j & \omega^{2j} & \dots & \omega^{j(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_j \\ \vdots \\ a_{n-1} \end{bmatrix}$$

FFT can be thought of as a clever way to choose points:

- Evaluations at many distinct points “collapse” together
- This is why we are left with $2T(n/2)$ work after division, instead of $4T(n/2)$ for a naive choice of points.

78/88

FFT-based Multiplication: Summary

- FFT works with 2^k points — Increases work by up to $2x$.
 - Product of two n^{th} -degree polynomial has degree $2n$
 - We need to work with $2n$ points, i.e., $4x$ increase in time.
- Requires inverting to coefficient representation after multiplication:

$$\vec{S}_n = M_n(\omega) \vec{A}_n$$

$$M_n^{-1}(\omega) \vec{S}_n = M_n^{-1}(\omega) M_n(\omega) \vec{A}_n = \vec{A}_n$$

It is easy to show that $M_n^{-1}(\omega) = M_n(-\omega)/n$, and hence:

$$\vec{A}_n * \vec{B}_n = \text{FFT}(\text{FFT}(\vec{A}_{2n}, \omega) \cdot \text{FFT}(\vec{B}_{2n}, \omega), \omega^{-1})/n$$

We are back to the convolution theorem!

79/88

More careful analysis ...

- Computations on complex or real numbers can lose precision.
 - For integer operations, we should work in some other ring — usually, we choose a ring based on modulo arithmetic.
 - Ex: in mod 33 arithmetic, 2 is the 10^{th} root of 1, i.e., $2^{10} \equiv 1 \pmod{33}$

More generally, 2 is the n^{th} root of unity modulo $(2^{n/2} + 1)$
- Point-wise additions and multiplications are not $O(1)$.
 - We are adding up to n numbers (“digits”) — we need $\Omega(\log n)$ bits
 - So, total cost increases by at least $\log n$, i.e., $O(n \log^2 n)$.
- [Schönhage-Strassen ’71] developed $O(n \log n \log \log n)$ algorithm: recursively apply their technique for “inner” operations.

80/88

Integer Multiplication Summary

- Algorithms implemented in libraries for arbitrary precision arithmetic, with applications in public key cryptography, computer algebra systems, etc.
- GNU MP is a popular library, uses various algorithms based on input size: naive, Karatsuba, Toom-3, Toom-4, or Schonhage-Strassen (at about 50K digits).
- Karatsuba is Toom-2. Toom-N is based on
 - Evaluating a polynomial at $2N$ points,
 - performing point-wise multiplication, and
 - interpolating to get back the polynomial, while
 - minimizing the operations needed for interpolation